

Getting and Cleaning Data

Ciprian Alexandru & Nicoleta Caragea

r-project.ro | @alexciopro

ICEA 2016 | Bucharest | ROMANIA | November 5th, 2016

Cleaning Data - Why?



How to clean the data?



"This is not what I meant when I said 'we need better data cleansing!'"

Topics

- ▶ Intro
- ▶ About R
 - ▶ R environment
 - ▶ R installation
 - ▶ GUI (RStudio and R Commander)
 - ▶ Packages
- ▶ Import and Export Data
- ▶ Data Manipulation
- ▶ Validate package
- ▶ Simputation package

Intro

Importing and cleaning data are the most important processes in data analysis. R is an efficient environment for detecting, diagnosing and finding data abnormalities. Along with the basic functions in R packages, there are packages dedicated to these processes. The `validate` package, a contribution of Mark van der Loo and Edwin de Jonge, help the data analysts to data validation process by checking data expectations about the data set. The `simputation` package aims to simplify missing value imputation using different methods like models and donor imputation. The models included in package are linear regression, robust linear regression, CART models and Random forest, respective the donor imputation methods k-nearest neighbour (based on gower's distance), sequential hotdeck (LOCF, NOCB), random hotdeck, predictive mean matching.

About R

What is R?

- ▶ R is a programming language and software environment for statistical computing and graphics
- ▶ The key point is the environment

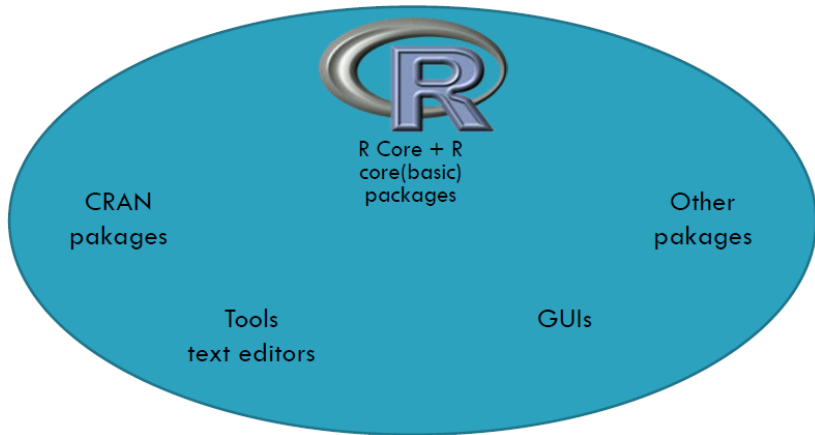
From where R comes?

- ▶ 1997 - Ross Ihaka and Robert Gentleman, professors of statistical at the Auckland University from New Zealand, starts to build a new software for statistical analysis and data graphical visualizations
- ▶ R is a dialect of S language (S was built by AT&T Bell Laboratories as a software for data analysis, statistical modeling, simulation and graphics)

Why R?

- ▶ R is supported by *academia*
- ▶ R is an *open source* initiative, similar with the Linux operating system or LaTeX markup language
- ▶ R is not just a statistics package, it's a *statistical programming language*
- ▶ R is designed to *overcome* the data scientist *problems*
- ▶ R is both *flexible*, powerful and endless

R Environment



R Quick installation

- ▶ Install R for UNIX platforms, Windows and MacOS from <https://www.r-project.org/>
- ▶ The Windows users just clicks, other users know better than others
- ▶ R version 3.3.2 (Sincere Pumpkin Patch) has been released on Monday 2016-10-31



The R Project for Statistical Computing

[Home]

Download

CRAN

R Project

About R

Logo

Contributors

What's New?

Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To **download R**, please choose your preferred [CRAN mirror](#).

If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

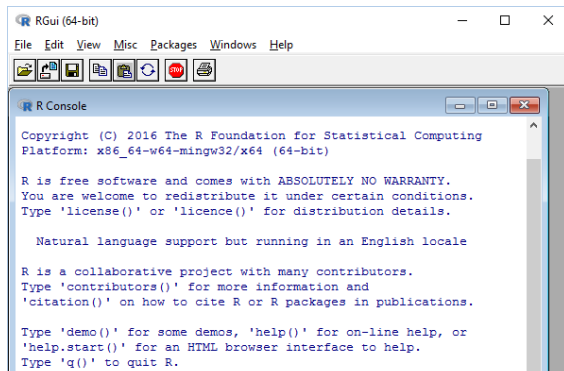
R-omania Team

GUI: R Console, R Studio, R Commander

- ▶ R Console: default GUI, include: the default multipledocumentinterface (MDI) and the single-document interface (SDI)
- ▶ R Studio: probably most complex GUI or integrated development environment (IDE) for R [<https://www.rstudio.com/products/rstudio/features/>, 2015-09-20]
- ▶ R Commander: contributed package Rcmdr: basic statistics GUI

R Console

- ▶ default GUI of R environment, included in R core
- ▶ > is the command prompt followed by a flashing cursor, meaning that R is waiting your reaction
- ▶ instructions/commands interpreted as functions



```
RGui (64-bit)
File Edit View Misc Packages Windows Help

Copyright (C) 2016 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.
```

Always useful. . .

- ▶ R is case sensitive: variable `ais` different from `A`
- ▶ Keywords: `if`, `else`, `repeat`, `while`, `function`, `for`, `in`, `next`, `break`, `TRUE`, `FALSE`, `NULL`, `Inf`, `NaN`, `NA`, `NA_integer_`, `NA_real_`, `NA_complex_`, and finally, `NA_character_`
- ▶ navigation commands executed: arrow Up and Down
- ▶ Ctrl+L clear the console content

R Studio

<https://www.rstudio.com/>



Take control of your R code

RStudio is an integrated development environment (IDE) for R. It includes a console, syntax-highlighting editor that supports direct code execution, as well as tools for plotting, history, debugging and workspace management. [Click here to see more RStudio features.](#)

RStudio is available in open source and commercial editions and runs on the desktop (Windows, Mac, and Linux) or in a browser connected to RStudio Server or RStudio Server Pro (Debian/Ubuntu, RedHat/CentOS, and SUSE Linux).



Desktop

Run RStudio on
your desktop

[RStudio
Desktop >](#)



Server

Centralize access
and computation

[RStudio Server >](#)

Announcing RStudio v1.0!

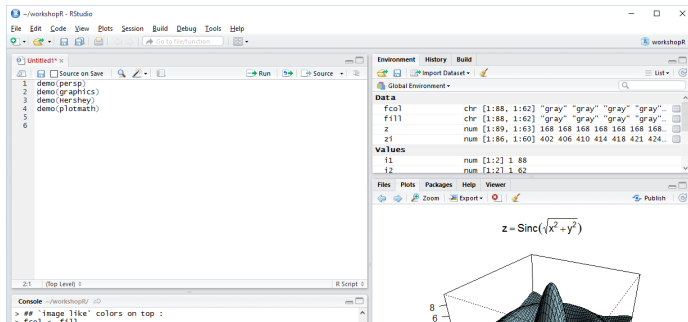
Today we're very pleased to announce the availability of RStudio Version 1.0! Version 1.0 is our 10th major release since the initial launch in February 2011 (see the full release history below), and our biggest ever! Highlights include:

- ▶ Authoring tools for R Notebooks.
- ▶ Integrated support for the sparklyr package (R interface to Spark).
- ▶ Performance profiling via integration with the profvis package.
- ▶ Enhanced data import tools based on the readr, readxl and haven packages.
- ▶ Authoring tools for R Markdown websites and the bookdown package.
- ▶ Many other miscellaneous enhancements and bug fixes.

R Studio IDE

4 working area:

- ▶ Text/Commands/Script editor
- ▶ Console
- ▶ Environment, History, Build
- ▶ Files, Plots, Packages, Help, Viewer



R Studio - features 1

- ▶ open source and commercial editions integrates the tools you use with R into a single environment
- ▶ available for Windows, Mac and Linux
- ▶ running on desktop, web browser and server
- ▶ efficient navigation to files and functions
- ▶ structure your work into projects
- ▶ integrated support for Git and subversion
- ▶ authoring HTML, PDF, Word Documents, and slide shows
- ▶ supports interactive graphics with Shiny and ggvis

R Studio - features 2

Integrated Development Environment (IDE):

- ▶ syntax highlighting
- ▶ code completion
- ▶ smart indentation
- ▶ execute R code directly from the source editor
- ▶ quickly jump to function definitions

Bring your workflow together:

- ▶ Integrated R help and documentation
- ▶ Easily manage multiple working directories using projects
- ▶ Workspace browser and data viewer

Authoring & Debugging:

- ▶ Interactive debugger to diagnose and fix errors quickly

R Commander


- ▶ `install.packages("Rcmdr")`
- ▶ `library(Rcmdr)`

Features:

- ▶ Data manipulation
- ▶ Statistics -basic statistical analyses
- ▶ Graphs - simple statistical graphs
- ▶ Models - numerical summaries, confidence intervals, hypothesis tests, diagnostics, and graphs for a statistical model, and for adding diagnostic quantities (eg, residuals) to the data set
- ▶ Distributions - probabilities, quantiles, and graphs of standard statistical distributions

R Commander (2)

- ▶ Script/editor window shows the command generated by user interactions with the menus.
- ▶ Edit/View data sets (reasonable small data sets)
- ▶ Data import from plain-text, Minitab, SPSS, or STATA
- ▶ Output window
- ▶ Messages window
- ▶ Graphics Device windows (appear separately)
- ▶ Save Graphs to different file type: bitmap, PDF, Postscript, EPS.
- ▶ Submit or Ctrl+rwill run your commands



R - Packages

Installed and loaded initially:

- ▶ stats; graphics; grDevices; utils; datasets; methods; base

Installed but not loaded:

- ▶ base; boot; class; cluster; codetools; compiler; datasets; foreign; graphics; grDevices; grid; KernSmooth; lattice; MASS; matrix; methods; mgcv; nlme; nnet; parallel; rpart; spatial; splines; stats; stats4; survival; tcltk; tools; utils

Contributed CRAN packages:

- ▶ ggplot2; zoo; ggmap; 7.168+

Other packages:

- ▶ not included in CRAN

Very used functions

- ▶ `ls()`
- ▶ `rm()`
- ▶ `install.packages("zoo")`
- ▶ `remove.packages("zoo")`
- ▶ `library(zoo)`

Operators & functions

- ▶ standard arithmetic operators: +, -, *, and /
- ▶ mathematical functions: sqrt, exp, and log
- ▶ relational operators <=, <, ==, >, >= and !=
- ▶ logical operators: | for OR and & for AND
- ▶ assignment operators: <- or = and ->

```
# Variable x gets value 2:
```

```
x <- 2
```

```
# Value 2 goes to variable x:
```

```
2 -> x
```


Operator syntax

- ▶ \$component extraction
- ▶ [[[indexing
- ▶ :sequence operator

```
x <- c(1:10)
x[(x < 5) | (x > 8)]
```

```
## [1] 1 2 3 4 9 10
```

Operator syntax (2)

```
1:5
```

```
## [1] 1 2 3 4 5
```

```
(a <-data.frame(name = c("Ion", "Maria"), income = c(1800,
```

```
##      name income
```

```
## 1   Ion   1800
```

```
## 2 Maria   2500
```

```
a$name
```

```
## [1] Ion   Maria
```

```
## Levels: Ion Maria
```

Operator syntax (3)

```
a[1]
```

```
##      name  
## 1     Ion  
## 2  Maria
```

```
a[2]
```

```
##    income  
## 1     1800  
## 2     2500
```

```
a[[1]]
```

```
## [1] Ion   Maria
```

Special values $+\text{Inf}$, $-\text{Inf}$, NaN

- ▶ R is properly infinite numerical values
- ▶ NaN-Not a Number
- ▶ Complex number:

```
sqrt(as.complex(-2))
```

```
## [1] 0+1.414214i
```

```
sqrt(-2+0i)
```

```
## [1] 0+1.414214i
```

Special values (2)

```
(a <- 2/0)
```

```
## [1] Inf
```

```
class(a)
```

```
## [1] "numeric"
```

```
exp(a)
```

```
## [1] Inf
```

```
exp(-a)
```

```
## [1] 0
```

Special values (3)

```
a - a
```

```
## [1] NaN
```

```
sqrt(a)
```

```
## [1] Inf
```

R objects

Five “atomic” classes of objects:

- ▶ character
- ▶ numeric (real numbers)
- ▶ integer
- ▶ complex
- ▶ logical (True/False)

R objects (2)

```
(x <- "a") # character
```

```
## [1] "a"
```

```
class(x)
```

```
## [1] "character"
```

```
(x <- 1) # numeric
```

```
## [1] 1
```

```
class(x)
```

```
## [1] "numeric"
```


R objects (3)

```
(x <- 1:5) # integer
```

```
## [1] 1 2 3 4 5
```

```
class(x)
```

```
## [1] "integer"
```

```
(x <- 2+3i) # complex
```

```
## [1] 2+3i
```

```
class(x)
```

```
## [1] "complex"
```

R objects (4)

```
(x <-TRUE) # logical
```

```
## [1] TRUE
```

```
class(x)
```

```
## [1] "logical"
```

R objects (5)

```
a <- 1  
b <- as.integer(1)  
a == b
```

```
## [1] TRUE
```

```
identical(a, b)
```

```
## [1] FALSE
```

R objects (6)

Near equality

```
(a <- 0.2 + 0.2 + 0.2)
```

```
## [1] 0.6
```

```
(b <- 0.6)
```

```
## [1] 0.6
```

```
a == b
```

```
## [1] FALSE
```

```
all.equal(a, b)
```

R objects (7)

```
a <- 1  
class(a)
```

```
## [1] "numeric"
```

```
typeof(a)
```

```
## [1] "double"
```

```
b <- 1:2  
class(b)
```

```
## [1] "integer"
```

R objects (8)

```
typeof(b)
```

```
## [1] "integer"
```

```
is.numeric(a)
```

```
## [1] TRUE
```

```
is.numeric(b)
```

```
## [1] TRUE
```

R - data structures

- ▶ factors
- ▶ atomic vector
- ▶ matrix
- ▶ array
- ▶ data frame
- ▶ list
- ▶ table

	Homogeneous	Heterogeneous
1d	Atomic vector	List
2d	Matrix	Data frame
nd	Array	

R - factor object

Factors - categorical data (unordered or ordered)

```
y <- c("yes", "no", "yes", "yes", "yes", "no")  
x <- c("yes", "no", "yes", "yes", "yes", "no")  
y <- as.factor(x)  
x
```

```
## [1] "yes" "no"  "yes" "yes" "yes" "no"
```

```
y
```

```
## [1] yes no  yes yes yes no  
## Levels: no yes
```


R - factor object (2)

```
str(x)
```

```
## chr [1:6] "yes" "no" "yes" "yes" "yes" "no"
```

```
str(y)
```

```
## Factor w/ 2 levels "no","yes": 2 1 2 2 2 1
```

R - factor object (3)

```
table(y)
```

```
## y  
##  no yes  
##   2  4
```

```
y
```

```
## [1] yes no  yes yes yes no  
## Levels: no yes
```

```
levels(y)
```

```
## [1] "no"  "yes"
```

R - factor object (3)

```
x <- factor(c("yes", "no", "yes", "yes", "no"), levels = c("yes", "no"))  
x
```

```
## [1] yes no  yes yes no
```


```
## Levels: yes no
```

R - data structures - data frame

- ▶ specific for data analysis/statisticians
- ▶ fundamental data structure by most of R's modeling software
- ▶ 2D vector (matrix), Heterogeneous data type
- ▶ list of vectors of equal length
- ▶ `data.frame(., row.names = NULL, check.rows = FALSE, check.names = TRUE, stringsAsFactors = default.stringsAsFactors())`
- ▶ `is.data.frame(x)`
- ▶ `dim(x)`
- ▶ `ncol(x)`
- ▶ `nrow(x)`
- ▶ `x[row, col]` or `x[observation, variable]`

R - Import and Export Data

The most flexible environment for data import



Data import (TXT)

```
read.table(file, header = FALSE, sep = "", quote = "\"", dec = ".",  
nrows = -1, skip = 0, colClasses = NA, ...)
```

```
date_txt <- read.table("pop_2015.txt", header=TRUE, sep=","  
head(date_txt)
```

```
##   varsta persoane  
## 1      0   191867  
## 2      1   193175  
## 3      2   180820  
## 4      3   185018  
## 5      4   206322  
## 6      5   214428
```

Data import (TXT) - 1

```
#fisier0 <- "https://raw.githubusercontent.com/alexcipro/i  
#date_txt_internet <- read.table(fisier0, header = TRUE, s  
#head(date_txt_internet)
```


Data import (CSV)

```
read.csv(file, header = TRUE, sep = ",", quote = "\"", dec = ".",  
fill = TRUE, comment.char = "", ...)
```

```
# read.csv
```

```
mydata <- read.csv("pop_2015.csv")
```

```
# similar cu read.table
```

```
mydata <- read.table("pop_2015.csv", head = TRUE, sep = ",",
```

Data import (Excel v1) - 1

- ▶ read first worksheet from mydata.xlsx
- ▶ first row contains variable names

```
#install.packages("xlsx")  
#install.packages("rJava")  
  
# work only in R-32 bit version  
library(rJava)  
library(xlsx)
```

```
## Loading required package: xlsxjars
```

Data import (Excel v1) - 2

```
mydata1 <- read.xlsx("mydata.xlsx", 1)  
head(mydata1)
```

##	Ozone	Solar.R	Wind	Temp	Month	Day
## 1	41	190	7.4	67	5	1
## 2	36	118	8.0	72	5	2
## 3	12	149	12.6	74	5	3
## 4	18	313	11.5	62	5	4
## 5	NA	NA	14.3	56	5	5
## 6	28	NA	14.9	66	5	6

```
# read data from the worksheet: Sheet1  
mydata2 <- read.xlsx("mydata.xlsx", sheetName = "Sheet1")  
head(mydata2)
```

Data import (Excel v2)

```
library(readxl)
mydata3 <- readxl::read_excel("mydata.xlsx")
head(mydata3)
```

```
## # A tibble: 6 × 6
##   Ozone Solar.R Wind Temp Month Day
##   <chr>   <chr> <dbl> <dbl> <dbl> <dbl>
## 1    41    190   7.4    67     5     1
## 2    36    118   8.0    72     5     2
## 3    12    149  12.6    74     5     3
## 4    18    313  11.5    62     5     4
## 5    NA     NA  14.3    56     5     5
## 6    28     NA  14.9    66     5     6
```

Data import (DBF)

```
require(foreign)
```

```
## Loading required package: foreign
```

```
mydata <- read.dbf("mydata.dbf")  
head(mydata)
```

##	OZONE	SOLAR.R	WIND	TEMP	MONTH	DAY
## 1	41	190	7	67	5	1
## 2	36	118	8	72	5	2
## 3	12	149	13	74	5	3
## 4	18	313	12	62	5	4
## 5	NA	NA	14	56	5	5
## 6	28	NA	15	66	5	6

Data import (SPSS)

```
library(foreign) # using the arguments  
use.value.labels don't convert value labels to factor  
levels mydata <- read.spss(file.choose(),  
use.value.labels = FALSE)
```

Data import (SAS)

```
# install.packages("Hmisc")  
# library(Hmisc)  
# mydata <- sasxport.get("d:/mydata.xpt")
```

Data import (stata)

```
# library(foreign)  
# mydata <- read.dta("mydata.dta")
```


Data export (TXT)

```
write.table(mydata2, "mydataw.txt", sep="\t")
```

Data export (Excel)

```
library(xlsx)  
write.xlsx(mydata2, "mydataaw.xlsx")
```

Data export (DBF)

```
library(foreign)  
write.dbf(mydata2, "mydataw.dbf")
```

SQL in R

```
#install.packages("sqldf")  
require(sqldf)
```

```
## Loading required package: sqldf
```

```
## Loading required package: gsubfn
```

```
## Loading required package: proto
```

```
## Loading required package: RSQLite
```

```
## Loading required package: DBI
```

```
myC02 <- C02  
head(C02, 3)
```

SQL - select all variables

```
# SQL  
s02 <- sqldf("select * from myCO2")
```

```
## Loading required package: tcltk
```

```
# R  
r02 <- myCO2[ , ]
```

SQL - select only one variable

```
# SQL  
s03 <- sqldf("select Type from myCO2")  
# R  
r03 <- myCO2[ , "Type"]
```

SQL - subset of variables

```
# SQL
s01 <- sqldf("select Type, conc from myCO2")
# R
r01 <- myCO2[, c("Type", "conc")]
# testing s01 vs. r01
all.equal(s01, r01)
```

```
## [1] TRUE
```

SQL - case sensitivity

```
# SQL is not case-sensitive  
# s04 <- sqldf("select type, coNC from myCO2")  
# R is case-sensitive  
# r04 <- myCO2[, c("type", "coNC")]
```


SQL - variable selection through number

```
head(myCO2[, c(1, 3, 5)], 3)
```

```
##   Plant Treatment uptake
## 1   Qn1 nonchilled   16.0
## 2   Qn1 nonchilled   30.4
## 3   Qn1 nonchilled   34.8
```

```
# the order of variables is important
head(myCO2[, c(5, 2)], 3)
```

```
##   uptake   Type
## 1   16.0 Quebec
## 2   30.4 Quebec
## 3   34.8 Quebec
```

SQL - variable selection through logic values

selection of variables/columns/fields by logic variables

```
head(myCO2[, c(TRUE, FALSE, FALSE, TRUE, FALSE)], 3)
```

```
##   Plant conc
## 1   Qn1   95
## 2   Qn1  175
## 3   Qn1  250
```

or

```
head(myCO2[, colnames(myCO2) > "d"], 3)
```

```
##   Plant   Type Treatment uptake
## 1   Qn1 Quebec nonchilled   16.0
## 2   Qn1 Quebec nonchilled   30.4
## 3   Qn1 Quebec nonchilled   24.8
```

SQL - selections by criteria (1)

```
# SQL
s05 <- sqldf("select * from myCO2 where uptake < 20")
# R
r05 <- myCO2[ myCO2[, "uptake"] < 20, ]
# or using with function
r05w <- with(myCO2, myCO2[uptake < 20, ]) # identical with
```

SQL - selections by criteria (2)

```
# SQL
```

```
s06 <- sqldf("select * from myCO2 where uptake < 20 and Type == 'Quebec'")
```

```
# R
```

```
r06 <- with(myCO2, myCO2[uptake < 20 & Type == 'Quebec', ])
```

SQL - first n observations

```
# SQL  
s07 <- sqldf("select * from myCO2 limit 6")  
# R  
r07 <- head(myCO2, 6)
```

SQL - NULL

```
r08 <- r06  
r08[2:4, 1] <- NA  
r08[5, 4] <- NA  
r08
```

##	Plant	Type	Treatment	conc	uptake
## 1	Qn1	Quebec	nonchilled	95	16.0
## 8	<NA>	Quebec	nonchilled	95	13.6
## 15	<NA>	Quebec	nonchilled	95	16.2
## 22	<NA>	Quebec	chilled	95	14.2
## 29	Qc2	Quebec	chilled	NA	9.3
## 36	Qc3	Quebec	chilled	95	15.1

SQL - Not NULL

```
# SQL  
s09 <- sqldf("select * from r08 where plant is not null")  
# R  
r09 <- with(r08, r08[!is.na(Plant), ])
```

SQL - is NULL

```
# SQL  
s10 <- sqldf("select * from r08 where plant is null")  
# R  
r10 <- with(r08, r08[is.na(Plant), ])
```


SQL - without missing values

```
# R  
na.omit(r08)
```

```
##      Plant   Type Treatment conc uptake  
## 1      Qn1 Quebec nonchilled   95    16.0  
## 36     Qc3 Quebec    chilled   95    15.1
```

R - Data manipulation

Data selection and manipulation (1)

- ▶ `which.max(x)`, `which.min(x)` - returns the index of the greatest/smallest element of `x`
- ▶ `rev(x)` - reverses the elements of `x`
- ▶ `sort(x)` - sorts the elements of `x` in increasing order; to sort in decreasing order: `rev(sort(x))`
- ▶ `cut(x,breaks)` - divides `x` into intervals (factors); `breaks` is the number of cut intervals or a vector of cut points
- ▶ `match(x,y)` returns a vector of the same length as `x` with the elements of `x` that are in `y` (NA otherwise)
- ▶ `which(x==a)` returns a vector of the indices of `x` if the comparison operation is true (TRUE), in this example the values of `i` for which `x[i] == a` (the argument of this function must be a variable of mode logical)

Data selection and manipulation (2)

- ▶ `choose(n,k)` computes the combinations of k events among n repetitions = $n! / [(n - k)!k!]$
- ▶ `na.omit(x)` suppresses the observations with missing data (NA)
- ▶ `na.fail(x)` returns an error message if x contains at least one NA
`complete.cases(x)` returns only observations (rows) with no NA
- ▶ `unique(x)` if x is a vector or a data frame, returns a similar object but with the duplicates suppressed
- ▶ `table(x)` returns a table with the numbers of the different values of x (typically for integers or factors)
- ▶ `split(x,f)` divides vector x into the groups based on f

Data selection and manipulation (3)

- ▶ `subset(x, ...)` returns a selection of `x` with respect to criteria (`...`, typically comparisons: `x$V1 < 10`); if `x` is a data frame, the option `select` gives variables to be kept (or dropped, using a minus)
- ▶ `na.fail(x)` returns an error message if `x` contains at least one NA
- ▶ `complete.cases(x)` returns only observations (rows) with no NA

Data reshaping (1)

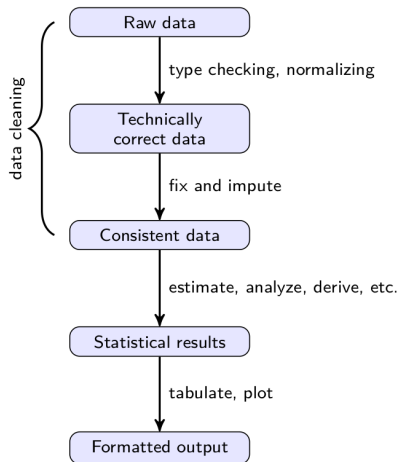
- ▶ `merge(a,b)` merge two data frames by common col or row names
- ▶ `stack(x,...)` transform data available as separate cols in a data frame or list into a single col
- ▶ `unstack(x,...)` inverse of `stack()`
- ▶ `rbind(...)`, `cbind(...)` combines supplied matrices, data frames, etc. by rows or cols
- ▶ `melt(data,id.vars,measure.vars)` changes an object into a suitable form for easy casting, (`reshape2` package)

Data reshaping (2)

- ▶ `cast(data, formula, fun)` applies fun to melted data using formula (reshape2 package)
- ▶ `recast(data, formula)` melts and casts in a single step (reshape2 package)
- ▶ `reshape(x, direction...)` reshapes data frame between 'wide' (repeated measurements in separate cols) and 'long' (repeated measurements in separate rows) format based on direction
- ▶ `aggregate(x, by, fun)` input df; output df; applies fun to subsets of x, as grouped based on index.

Statistical analysis steps

Concepts



Technically correct data

- ▶ Well-defined format (data structure)
- ▶ Well-defined types (numbers, date/time, string, categorical...)
- ▶ Statistical units can be identified (persons, transactions, phone calls...)
- ▶ Variables can be identified as properties of statistical units.
- ▶ Note: tidy data \subset technically correct data

Consistent data

- ▶ Data satisfies demands from domain knowledge (validation process)

Dirty tabular data

`read.table` vs. `readr::read_csv`

- ▶ `read.table`: R's swiss army knife
 - ▶ fairly strict (no sniffing)
 - ▶ very flexible
 - ▶ interface could be cleaner
- ▶ `readr::read_csv`
 - ▶ easy to switch between strict/lenient parsing
 - ▶ compact control over column types
 - ▶ fast
 - ▶ clear reports of parsing failure

reading with read.table (1)

```
dat <- read.table(file = "table/unnamed.csv",  
  header = FALSE,  
  col.names = c("age", "height"),  
  stringsAsFactors = FALSE,  
  sep = ",")
```

dat

```
##    age height  
## 1   21    6.0  
## 2   42    5.9  
## 3   18    5.7*  
## 4   21   <NA>
```

reading with read.table (2)

```
class(dat$height)
```

```
## [1] "character"
```

```
dat$height <- as.numeric(dat$height)
```

```
## Warning: NAs introduced by coercion
```

```
dat
```

```
##   age height
## 1  21    6.0
## 2  42    5.9
## 3  18     NA
## 4  21     NA
```

define colClasses

- ▶ this will generate an error

```
#dat <- read.table(  
# file = "table/unnamed.csv",  
# header = FALSE,  
# col.names = c("age", "height"),  
# colClasses = c("numeric", "numeric"),  
# stringsAsFactors = FALSE,  
# sep = ",")
```

reading with the readr package (1)

- parse columns as 'number' (flexible)

```
readr::read_csv("table/unnamed.csv",  
  col_names=c("age", "height"),  
  col_types="nn")
```

```
## # A tibble: 4 × 2  
##   age height  
##   <dbl> <dbl>  
## 1    21    6.0  
## 2    42    5.9  
## 3    18    5.7  
## 4    21    NA
```


reading with the readr package (2)

- parse columns as 'double' (strict)

```
readr::read_csv("table/unnamed.csv",  
  col_names=c("age", "height"),  
  col_types="dd")
```

```
## Warning: 1 parsing failure.
```

```
## row      col                expected actual  
##    3 height no trailing characters      *
```

```
## # A tibble: 4 × 2
```

```
##   age height
```

```
##   <dbl> <dbl>
```

```
## 1    21    6.0
```

```
## 2    42    5.9
```

Real dirty data (1)

```
source("parse/parse_outfile.R")  
to <- read_tof("parse/skylark-1d.out")  
to
```

```
## TRIM 3.61 : TRend analysis and Indices for Monitoring  
## STATISTICS NETHERLANDS  
##  
## Date/Time: 4-7-2016 15:08:28  
##  
## Title : skylark-1d  
##  
## Comment: Example 1; using linear trend model  
##  
## The following 5 variables have been read from file:  
## F:\TRIM\TRIM manual demo\skylark dat
```

Real dirty data (2)

```
get_n_site(to)
```

```
## [1] 55
```

```
get_n_site_stringr(to)
```

```
## [1] 55
```

```
get_time_indices(to)
```

```
##      Time  Model std.err. Imputed std.err..1  
## 1      1 1.0000  1.0000      NA      NA  
## 2      2 1.0496  0.0149  0.8948  0.0410  
## 3      3 1.1017  0.0312  0.9777  0.0601  
## 4      4 1.1562  0.0401  0.9700  0.0670
```

Lessons learned

- ▶ (base) R has great text processing tools.
- ▶ Need to work with regular expressions¹
- ▶ Write many small functions extracting single data elements.
- ▶ Don't overgeneralize: adapt functions as you meet new input.
- ▶ Smart use of existing tools (`read.table(text=)`)

¹Mastering Regular Expressions (2006) by Jeffrey Friedl is a great resource

Packages for standard format parsing

- ▶ `jsonlite`: parse JSON files
- ▶ `yaml`: parse yaml files
- ▶ `xml2`: parse XML files
- ▶ `rvest`: scrape and parse HTML files

String normalization

Bring a text string in a standard format, e.g.

- ▶ Standardize upper/lower case (casefolding)
 - ▶ stringr: `str_to_lower`, `str_to_upper`, `str_to_title`
 - ▶ base R: `tolower`, `toupper`
- ▶ Remove accents (transliteration)
 - ▶ stringi: `stri_trans_general`
 - ▶ base R: `iconv`
- ▶ Re-encoding
 - ▶ stringi: `stri_encode`
 - ▶ base R: `iconv`
- ▶ Uniformize encoding (unicode normalization)
 - ▶ stringi: `stri_trans_nfkc` (and more)

Approximate text matching: edit-based distances

Distance	Allowed operation			
	substitution	deletion	insertion	transposition
Hamming	✓	✗	✗	✗
LCS	✗	✓	✓	✗
Levenshtein	✓	✓	✓	✗
OSA	✓	✓	✓	✓*
Damerau-Levenshtein	✓	✓	✓	✓

*Substrings may be edited only once.

"leela" → "leea" → "leia"

```
stringdist::stringdist("leela","leia",method="dl")
```

```
## [1] 2
```

Some pointers for approximate matching

- ▶ Normalisation and approximate matching are complementary
- ▶ See Mark Van Der Loo [useR2014 talk](#) or [paper](#) on stringdist for more distances
- ▶ The [fuzzyjoin](#) package allows fuzzy joining of datasets

Other good stuff

- ▶ lubridate: extract dates from strings

```
lubridate::dmy("17 December 2015")
```

```
## [1] "2015-12-17"
```

- ▶ tidyr: many data cleaning operations to make your life easier
- ▶ readr: Parse numbers from text strings

```
readr::parse_number(c("2%", "6%", "0.3%"))
```

```
## [1] 2.0 6.0 0.3
```

Validation and Imputation

The `validate` package, in summary

- ▶ Make data validation rules explicit
- ▶ Treat them as objects of computation
 - ▶ store to / read from file
 - ▶ manipulate
 - ▶ annotate
- ▶ Confront data with rules
- ▶ Analyze/visualize the results

Use rules to correct data

Main idea

Rules restrict the data. Sometimes this is enough to derive a correct value uniquely.

Examples

- ▶ Correct typos in values under linear restrictions
 - ▶ $123 + 45 \neq 177$, but $123 + \underline{54} = 177$.
- ▶ Derive imputations from values under linear restrictions
 - ▶ $123 + \text{NA} = 177$, compute $177 - 123 = 54$.

Both can be generalized to systems $\mathbf{Ax} \leq \mathbf{b}$.

Validate

```
library(magrittr)
library(validate)
data(retailers)
head(retailers, 3)
```

```
##   size incl.prob staff turnover other.rev total.rev sta
## 1  sc0      0.02   75      NA      NA      1130
## 2  sc3      0.14    9    1607      NA      1607
## 3  sc3      0.14   NA    6886     -33      6919
##   total.costs profit vat
## 1      18915  20045  NA
## 2       1544    63  NA
## 3       6493   426  NA
```

A first glance

```
retailers %>%  
  check_that(other.rev > 0, profit < turnover) %>%  
  summary()
```

```
##    rule items passes fails nNA error warning      expre  
## 1   V1    60     23     1  36 FALSE   FALSE    other.re  
## 2   V2    60     48     4   8 FALSE   FALSE profit < tur
```

Define rules for reuse (1)

```
v <- validator(staff >= 0,  
  turnover >= 0,  
  other.rev >= 0,  
  total.rev >= 0,  
  turnover + other.rev == total.rev,  
  if (staff > 0) staff.costs > 0  
)
```

Define rules for reuse (2)

```
v
```

```
## Object of class 'validator' with 6 elements:  
## V1: staff >= 0  
## V2: turnover >= 0  
## V3: other.rev >= 0  
## V4: total.rev >= 0  
## V5: turnover + other.rev == total.rev  
## V6: !(staff > 0) | staff.costs > 0
```

```
summary(v)
```

```
##   block nvar rules linear  
## 1      1     2     2      1  
## 2      2     3     4      4
```


getters and setters for rule metadata (1)

```
created(v)
```

```
## [1] "2016-11-05 12:58:44 EET" "2016-11-05 12:58:44 EET"  
## [3] "2016-11-05 12:58:44 EET" "2016-11-05 12:58:44 EET"  
## [5] "2016-11-05 12:58:44 EET" "2016-11-05 12:58:44 EET"
```

```
origin(v)
```

```
## [1] "command-line" "command-line" "command-line" "command-line"  
## [5] "command-line" "command-line"
```

```
names(v)
```

```
## [1] "V1" "V2" "V3" "V4" "V5" "V6"
```

getters and setters for rule metadata (2)

```
description(v)
```

```
## [1] "" "" "" "" "" ""
```

```
cf <- confront(retailers, v)  
cf
```

```
## Object of class 'validation'  
## Call:  
##      confront(x = retailers, dat = v)  
##  
## Confrontations: 6  
## With fails      : 2  
## Warnings        : 0  
## Errors          : 0
```

getters and setters for rule metadata (3)

```
summary(cf)
```

##	rule	items	passes	fails	nNA	error	warning
## 1	V1	60	54	0	6	FALSE	FALSE
## 2	V2	60	56	0	4	FALSE	FALSE
## 3	V3	60	23	1	36	FALSE	FALSE
## 4	V4	60	58	0	2	FALSE	FALSE
## 5	V5	60	19	4	37	FALSE	FALSE
## 6	V6	60	50	0	10	FALSE	FALSE

```
##                                expression
## 1                                staff >= 0
## 2                                turnover >= 0
## 3                                other.rev >= 0
## 4                                total.rev >= 0
```

R-omania Team

getters and setters for rule metadata (4)

```
aggregate(cf, by="record") %>% head(3)
```

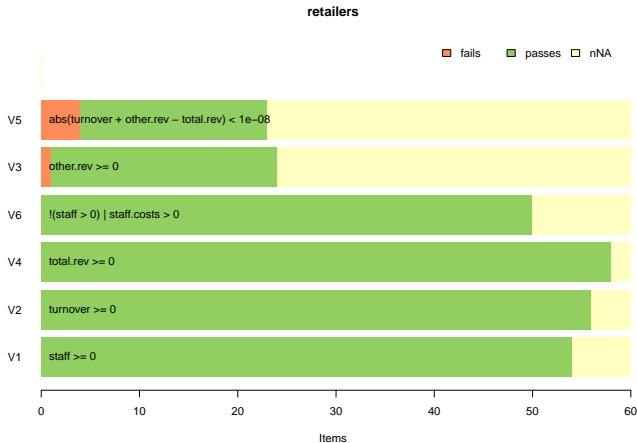
##	npass	nfail	nNA	rel.pass	rel.fail	rel.NA
## 1	2	0	4	0.3333333	0.0000000	0.6666667
## 2	4	0	2	0.6666667	0.0000000	0.3333333
## 3	3	2	1	0.5000000	0.3333333	0.1666667

```
sort(cf, by="rule") %>% head(3)
```

##	npass	nfail	nNA	rel.pass	rel.fail	rel.NA
## V5	19	4	37	0.3166667	0.06666667	0.6166667
## V3	23	1	36	0.3833333	0.01666667	0.6000000
## V6	50	0	10	0.8333333	0.00000000	0.1666667

getters and setters for rule metadata (5)

```
barplot(cf, main="retailers")
```



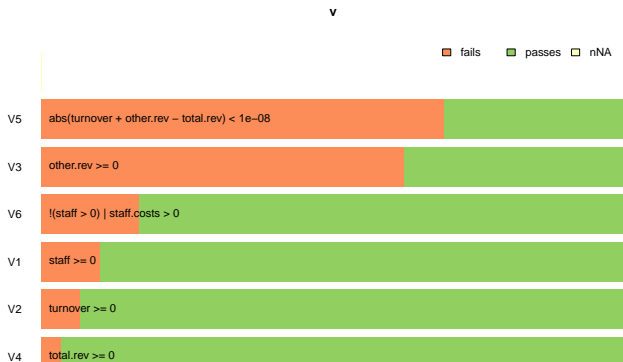
setting options (1)

```
retailers %>%  
  confront(v, lin.eq.eps=1e-8) %>%  
  barplot()
```



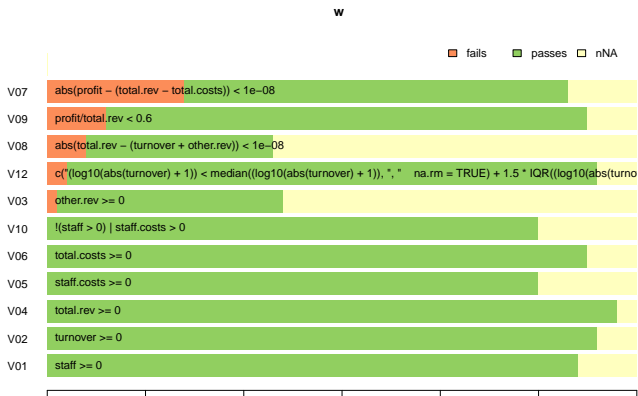
setting options (2)

```
retailers %>%  
  confront(v, na.value=FALSE) %>%  
  barplot()
```



Reading from file

```
w <- validator(.file="validate/rules.R")  
confront(retailers,w) %>% barplot()
```



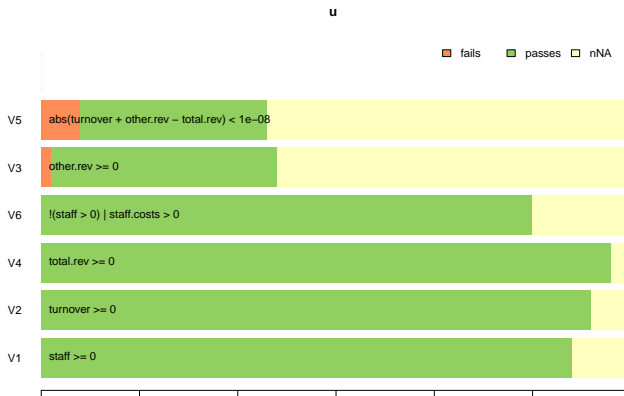
Rich metadata: yaml files (1)

```
u <- validator(.file="validate/rules.yaml")  
u
```

```
## Object of class 'validator' with 6 elements:  
##  V1 [nonnegative staff]           : staff >= 0  
##  V2 [nonnegative turnover]        : turnover >= 0  
##  V3 [nonnegative other rev]       : other.rev >= 0  
##  V4 [nonnegative total rev]       : total.rev >= 0  
##  V5 [Revenue balance]             : turnover + other  
##  V6 [positive staff => positive costs]: !(staff > 0) | s
```

Rich metadata: yaml files (2)

```
confront(retailers,u) %>%  
  barplot()
```



Error localization

Notes on `errorlocate`

- ▶ For in-record rules
- ▶ Support for
 - ▶ linear (in)equality rules
 - ▶ Conditionals on categorical variables (if male then not pregnant)
 - ▶ Mixed conditionals (has job then age ≥ 15)
 - ▶ Conditionals w/linear predicates (staff > 0 then staff cost > 0)
- ▶ Optimization is mapped to MIP problem.

Missing values

Mechanisms (Rubin):

- ▶ **MCAR**: missing completely at random
- ▶ **MAR**: $P(Y = \text{NA})$ depends on value of X
- ▶ **MNAR**: $P(Y = \text{NA})$ depends on value of Y

Imputation

Purpose of imputation vs prediction

- ▶ Prediction: estimate a single value (often for a single use)
- ▶ Imputation: estimate values such that the completed data set allows for valid inference^a

^aThis is very difficult!

Imputation methods

- ▶ Deductive imputation
- ▶ Imputation based on predictive models
- ▶ Donor imputation (knn, pmm, sequential/random hot deck)

Predictive model-based imputation

$$\hat{y} = \hat{f}(\mathbf{x}) + \epsilon$$

e.g. Linear regression

$$\hat{y} = \alpha + \mathbf{x}^T \hat{\beta} + \epsilon$$

- ▶ Residual:
 - ▶ $\epsilon = 0$ Impute expected value
 - ▶ ϵ drawn from observed residuals e
 - ▶ $\epsilon \sim N(0, \sigma)$ parametric residual, $\hat{\sigma}^2 = \text{var}(e)$
- ▶ Multiple imputation (Bayesian bootstrap)
 - ▶ Draw β from parametric distribution, impute multiple times.

Donor imputation (hot deck)

Method variants:

- ▶ **Random hot deck:** copy value from random record.
- ▶ **Sequential hot deck:** copy value from previous record.
- ▶ **k -nearest neighbours:** draw donor from k nearest neighbours
- ▶ **Predictive mean matching:** copy value closest to prediction

Donor pool variants:

- ▶ per variable
- ▶ per missing data pattern
- ▶ per record

Note on multivariate donor imputation

Many multivariate methods seem relatively *ad hoc*, and more theoretical and empirical comparisons with alternative approaches would be of interest.

Andridge and Little (2010) *A Review of Hot Deck Imputation for Survey Non-response*. Int. Stat. Rev. **78**(1) 40-64

Methods supported by `simputation`

- ▶ Model based (optionally add [non-]parametric random residual)
 - ▶ linear regression
 - ▶ robust linear regression
 - ▶ CART models
 - ▶ Random forest
- ▶ Donor imputation (including various donor pool specifications)
 - ▶ k-nearest neighbour (based on `gower`'s distance)
 - ▶ sequential hotdeck (LOCF, NOCB)
 - ▶ random hotdeck
 - ▶ Predictive mean matching
- ▶ Other
 - ▶ (groupwise) median imputation (optional random residual)
 - ▶ Proxy imputation (copy from other variable)

Simputation package

Investigate missing data patterns

```
library(VIM)
```

```
## Loading required package: colorspace
```

```
## Loading required package: grid
```

```
## Loading required package: data.table
```

```
## VIM is ready to use.
```

```
## Since version 4.0.0 the GUI is in its own package VIMGUI
```

```
##
```

```
##
```


```
    Please use the package to use the new (and old)
```

```
## Suggestions and bug-reports can be submitted at: https://
```

```
##
```

plot missing data patterns (1)


```
a <- VIM::aggr(retailers[3:9], sortComb=TRUE, sortVar=TRUE)
```



plot missing data patterns (2)


#

```
VIM::pbox(retailers[3:9],las=2)
```



plot missing data patterns (3)

```
dat <- log10(abs(retailers[c(3,5)]))  
VIM::marginplot(dat, las=1, pch=16)
```



plot missing data patterns (4)

```
# testing means
t.test(log(staff) ~ is.na(other.rev), data=retailers)

##
##  Welch Two Sample t-test
##
## data:  log(staff) by is.na(other.rev)
## t = 2.7464, df = 46.014, p-value = 0.008572
## alternative hypothesis: true difference in means is not
## 95 percent confidence interval:
##  0.1985149 1.2880867
## sample estimates:
## mean in group FALSE  mean in group TRUE
##           2.329996           1.586695
```

Impute values using simulation

Linear model to impute three variables:

```
d1 <- impute_lm(retailers, turnover ~ staff)
validate::cells(retailers, d1)
```

```
## Object of class cellComparison:
##
##      validate::cells(retailers, d1)
##
##              D0001 D0002
## cells              600    600
## available          520    523
## missing             80     77
## still_available    520    520
## unadapted          520    520
## adapted             0      0
## imputed            0      2
```

Use staff as predictor for multiple variables

```
d2 <- impute_lm(retailers, turnover + other.rev + total.rev  
validate::cells(retailers, d2)
```

```
## Object of class cellComparison:  
##  
##      validate::cells(retailers, d2)  
##  
##              D0001 D0002  
## cells              600    600  
## available          520    559  
## missing            80     41  
## still_available    520    520  
## unadapted          520    520  
## adapted            0      0  
## imputed            0      0
```


Impute everything excepts staff as a function of staff (1)

```
d3 <- impute_lm(retailers, . ~ staff ~ staff)
validate::cells(retailers, d3)
```

```
## Object of class cellComparison:
##
##      validate::cells(retailers, d3)
##
##               D0001 D0002
## cells           600    600
## available       520    588
## missing         80     12
## still_available 520    520
## unadapted       520    520
## adapted         0      0
## imputed         0      68
```


Impute everything excepts staff as a function of staff (2)

`aggr(d3)`



Chain methods

```
d4 <- retailers %>%  
  impute_lm(.-staff ~ staff) %>% # linear model  
  impute_median(. ~ size)        # group median  
aggr(d4)
```



Some other methods

copy value from proxy

```
d5 <- impute_proxy(retailers, total.rev ~ vat)
validate::cells(retailers,d5)
```

```
## Object of class cellComparison:
##
##      validate::cells(retailers, d5)
##
##              D0001 D0002
## cells              600    600
## available          520    521
## missing             80     79
## still_available    520    520
## unadapted          520    520
## adapted              0      0
## imputed              0      1
```

CART model imputation

```
d6 <- impute_cart(retailers, total.rev ~ .)
```


robust linear model with parametric residuals added

```
d7 <- impute_rlm(retailers,  
  total.rev ~ staff,  
  add_residual = "normal")
```

variance of estimation, including imputation by bootstrap (1)

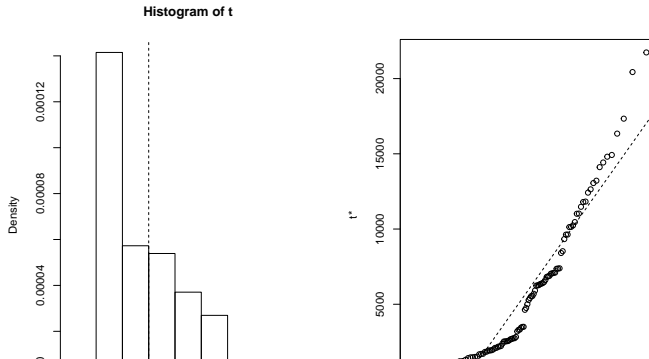
```
stat <- function(dat,i){  
  dat <- dat[i,,drop=FALSE]  
  dat %<>% impute_lm(staff.costs ~ staff + total.rev) %>%  
    impute_lm(staff.costs ~ staff) %>%  
    impute_median(staff.costs ~ size) %>%  
    impute_const(staff.costs ~ 0)  
  mean(dat[, "staff.costs"], na.rm=TRUE)  
}  
stat(retailers, seq_len(nrow(retailers)))
```

```
## [1] 6398.378
```

```
library(boot)
```

variance of estimation, including imputation by bootstrap (2)

```
b = boot(data=retailers, statistic = stat,R=100)  
plot(b)
```



Credits

- ▶ `deductive` Mark van der Loo, Edwin de Jonge
- ▶ `errorlocate` Edwin de Jonge, Mark van der Loo
- ▶ `gower` Mark van der Loo
- ▶ `jsonlite` Jeroen Ooms, Duncan Temple Lang, Lloyd Hilaiel
- ▶ `magrittr` Stefan Milton Bache, Hadley Wickham
- ▶ `rex` Kevin Ushey, Jim Hester, Robert Krzyzanowski
- ▶ `simputation` Mark van der Loo
- ▶ `stringdist` Mark van der Loo, Jan van der Laan, R Core, Nick Logan
- ▶ `stringi` Marek Gagolewski, Bartek Tartanus
- ▶ `stringr` Hadley Wickham, RStudio
- ▶ `tidyr` Hadley Wickham, RStudio
- ▶ `validate` Mark van der Loo, Edwin de Jonge
- ▶ `VIM` Matthias Templ, Andreas Alfons, Alexander Kowarik, Bernd Prantner
- ▶ `xm12` Hadley Wickham, Jim Hester, Jeroen Ooms, RStudio, R foundation

Thank you for your kind attention!



Learn R - Invata R

Link:

[<http://www.r-project.ro/invatar/intro/index.html>]

Release date: 05.nov.2016

← → ↻ ⓘ www.r-project.ro/invatar/intro/index.html ☆

Invata R

INVATA R IN 10 MINUTE PE ZI

🏠 Intro Despre ImpExp Functii Pachete

R Intro

- Introducere
- Ce este R?
- De ce R?
- Caracteristici R
- Consola R
- Diverse comenzi
- Configurarea mediului R
- Intotdeauna utile...
- Uneori utile...
- Tot despre mediul R
- Exemple 1
- Exemple 2
- Mediul R - personalizare
- Mediul R

Invata R Introducere

Acesta este primul pas pe care il puteti face in utilizarea mediului R.

Parcurgeti materialul prezentat si efectuati exemplele puse la dispozitie pentru a va deprinde cu experimentele.

Autor: Ciprian Alexandru

eContact: alexcipro @ Yahoo | Gmail | GitHub | Twitter

R contributions: [<http://www.r-project.ro/>]

R-omania Team

Evenimente
Cercetare
Echipa

R links

R-project.org
RStudio.com
R-bloggers.com

ia Team